

# AirBnb listings file modeling - final regressions

This file loads the cleaned data, runs it through the pipeline and runs the regressors on the data with all the hyperparameters tuned and ready.

```
In [40]: 1 # Put these at the top of every notebook, to get automatic reloading and inline
2 from IPython.core.display import display, HTML
3 import pandas as pd
4 import numpy as np
5 import warnings
6 import ast
7 warnings.filterwarnings('ignore')
8
9 %reload_ext autoreload
10 %autoreload 1
11 %matplotlib inline
12
13 pd.set_option('display.max_rows', 500)
14 pd.set_option('display.max_columns', 500)
15 pd.set_option('display.width', 1000)
16
17 display(HTML("<style>.container { width:100% !important; }</style>"))
```

```
In [41]: 1 import os
2 import seaborn as sns
3 import pandas as pd
4 import math
5
6 import sklearn.model_selection as cv
7
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.model_selection import train_test_split
10 from sklearn.decomposition import PCA
11 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
12 from sklearn.tree import DecisionTreeRegressor
13 from sklearn.linear_model import Ridge, Lasso
14 from sklearn.model_selection import GridSearchCV
15
16 from sklearn.metrics import mean_squared_error as MSE
17
18 from imblearn.over_sampling import SMOTE
19
20 from Utils.UtilsViz import *
21 from Utils.DataUtils import *
22
23 RANDOM_SEED = 42
```

```
In [42]: 1 sns.set(font_scale=2, style="whitegrid")
```

```
In [43]: 1 data_path = os.path.join(os.getcwd(), "../data/cleaned_listings.csv")
2 listings = pd.read_csv(data_path, index_col="id")
3 display(listings.shape)
```

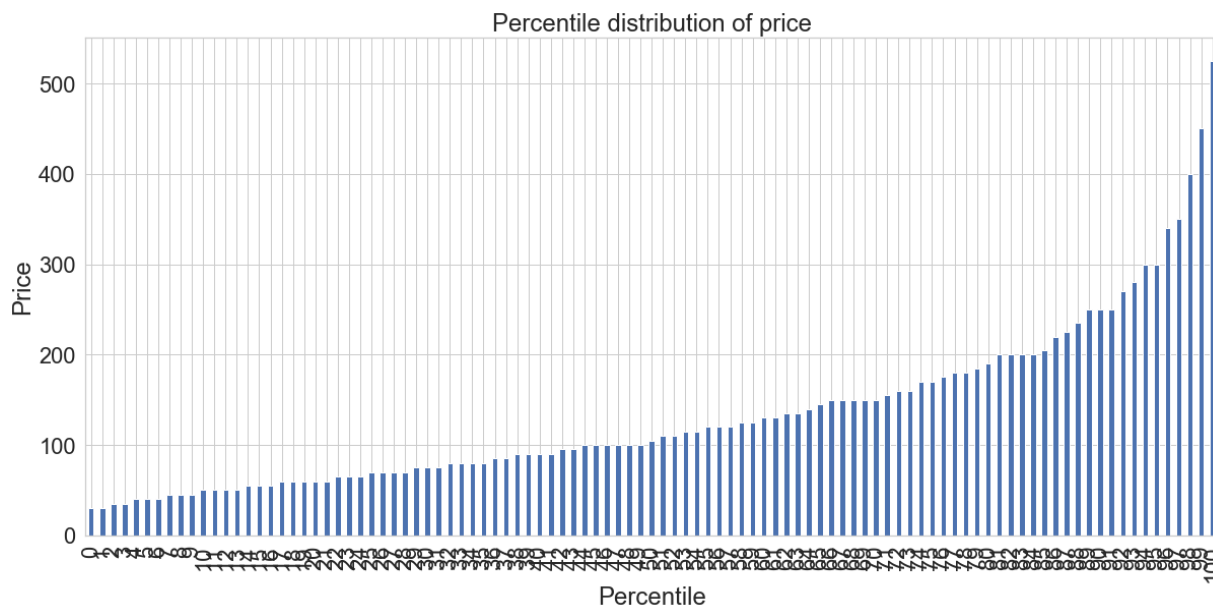
(48855, 65)

## Plot the dstribution

Let's plot the percentile for price

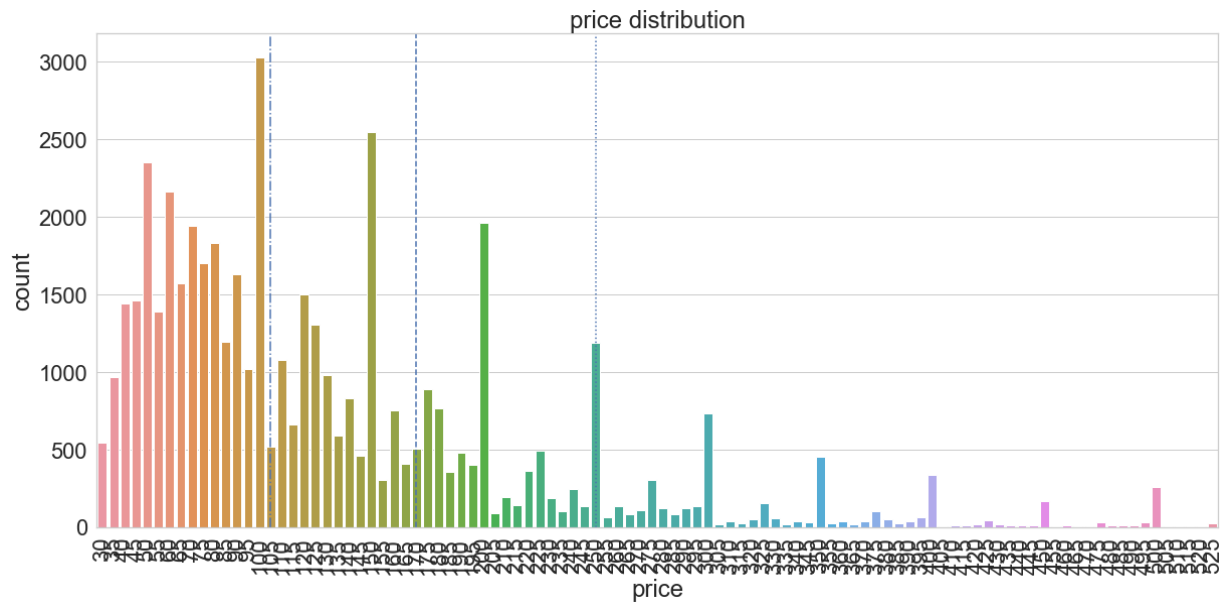
```
In [44]: 1 percentiles = list(range(0,101, 1))
2 price_percentile = {}
3 for p in percentiles:
4     price_percentile[p] = np.percentile(listings['price'].values, p)
5
6 price_percentile = pd.DataFrame.from_dict(price_percentile, orient='index')
7 price_percentile.plot(kind='bar', figsize=(20,9), grid=True, legend=False)
8 plt.title("Percentile distribution of price")
9 plt.xlabel("Percentile")
10 plt.ylabel("Price")
```

Out[44]: Text(0, 0.5, 'Price')



```
In [45]: 1 f, ax = plt.subplots(1,1,figsize=(20,9))
2         g = sns.countplot(x="price", data=listings, ax=ax)
3         t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
4         t = g.set_title("price distribution")
5         median_idx = np.where(np.sort(listings["price"].unique())==listings["price"]
6                                .median())[0][0]
7         plt.axvline(x=median_idx, linestyle="-.")
8         percentile_75_idx = np.where(np.sort(listings["price"].unique())==price_perc
9                                       [0.75])[0][0]
10        plt.axvline(x=percentile_75_idx, linestyle="--")
11        percentile_90_idx = np.where(np.sort(listings["price"].unique())==price_perc
12                                      [0.9])[0][0]
13        plt.axvline(x=percentile_90_idx, linestyle=":")
```

Out[45]: <matplotlib.lines.Line2D at 0xf46fc50>



Quick helper functions

```
In [46]: 1 def roundto(row, base=5):
2         return int(base * round(float(row) / base))
3
4         # Get the index of the price columns
5         def get_index(vallist, val):
6             return vallist.index(val)
```

## Oversampling using SMOTE

```
In [47]: 1 def check_rep(row):
2         if (row <= 200) | (row==250) | (row==350) | (row==450) | (row==550) :
3             return 0
4         elif (row > 200) & (row < 300) & (row != 250):
5             return 1
6         elif (row > 300) & (row < 400) & (row != 350):
7             return 2
8         else:
9             return 3
10
11 listings["flag_ur"] = listings["price"].apply(check_rep)
```

```
In [48]: 1 vcs = listings["flag_ur"].value_counts()
2         vcs
```

```
Out[48]: 0    43321
1      3093
3      1624
2       817
Name: flag_ur, dtype: int64
```


```
In [49]: 1 ycol = ["flag_ur"]
2         xcol = [i for i in listings.columns if i not in ycol]
3
4         x = listings[xcol].values
5         y = listings[ycol].values
6
7         smote_sampling_strategy = {
8             1: int(vcs[1]*2)
9             ,2: int(vcs[2]*2)
10            ,3: int(vcs[3]*2)
11        }
12         sm = SMOTE(random_state=RANDOM_SEED, sampling_strategy=smote_sampling_strate
13         # Fit the smote onto the sample
14         x_new, y_new = sm.fit_sample(x, y)
15
16         # Drop the flag column
17         listings.drop(labels=["flag_ur"], axis=1, inplace=True)
18
19         # -----
20         # Overwrite X and Y
21         price_index = get_index(list(listings.columns), "price")
22
23         y = x_new[:, price_index]
24         x = np.delete(x_new, price_index, axis=1)
25         for i in range(len(y)):
26             y[i] = roundto(y[i])
```

```
In [50]: 1 print(
2         " Old size :", listings.shape, "\n",
3         "New size :", x.shape
4         )
```

```
Old size : (48855, 65)
New size : (54389, 64)
```

```
In [51]: 1 f, ax = plt.subplots(1, 1, figsize=(15, 9), sharey=True)
2         g2 = sns.distplot(y, ax=ax, kde=False, label="SMOTE dist.")
3         g1 = sns.distplot(listings["price"], ax=ax, kde=False, label="Original dist.")
4         t = g1.set(title="Price distribution", ylabel="count")
5         plt.legend()
```

Out[51]: <matplotlib.legend.Legend at 0xf5468d0>



```
In [52]: 1 x_cols = listings.drop(['price'], axis=1)
2         X = pd.DataFrame(data=x, columns=x_cols.columns)
```

## Train test split

```
In [53]: 1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, ra
```

## Standardisation

```
In [54]: 1 standard_scaler = StandardScaler()
2         x_train = standard_scaler.fit_transform(x_train)
3         x_test = standard_scaler.transform(x_test)
```

## Prediction

```
In [55]: 1 # Quick placeholder variable
2         rmse_dict = {}
3         rmse_m_dict = {}
```

## a. Ridge Regression

```
In [73]: 1 rid = Ridge(alpha=1e-08)
2         rid.fit(X=x_train, y=y_train)
3         y_pred_train = rid.predict(X=x_train)
4         y_pred_test = rid.predict(X=x_test)
5
6         mse_train = MSE(y_train, y_pred_train)
7         mse_test = MSE(y_test, y_pred_test)
8
9         rmse_train = mse_train**(1/2)
10        rmse_test = mse_test**(1/2)
11
12        rmse_dict["Ridge"] = {"Train":rmse_train,"Test":rmse_test}
13
14        print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
15        print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 64.40

Test set RMSE Scaled: 65.17


What would be the RMSE for the majority (80 percentile) of the data?

```
In [74]: 1 th = price_percentile.iloc[80, :].values[0]
2
3         tmpdf = pd.DataFrame({"y_train":y_train, "y_pred_train":y_pred_train})
4         tmpdf = tmpdf[tmpdf["y_train"] <= th]
5
6         tmpdf2 = pd.DataFrame({"y_test":y_test, "y_pred_test":y_pred_test})
7         tmpdf2 = tmpdf2[tmpdf2["y_test"] <= th]
8
9         mse_train = MSE(tmpdf["y_train"].values, tmpdf["y_pred_train"].values)
10        mse_test = MSE(tmpdf2["y_test"].values, tmpdf2["y_pred_test"].values)
11
12        rmse_train = mse_train**(1/2)
13        rmse_test = mse_test**(1/2)
14
15        rmse_m_dict["Ridge"] = {"Train":rmse_train,"Test":rmse_test}
16
17        print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
18        print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 47.92

Test set RMSE Scaled: 49.45

```
In [75]: 1 f, ax = plt.subplots(1,1, figsize=(30, 9), sharex=False)
2 g = sns.pointplot(x=y_train.astype(int), y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test.astype(int), y=y_pred_test, ax=ax, color="red")
4 t = g.set_xlabel("Price")
5 t = g.set_ylabel("Predicted price distribution")
6 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
7 t = g.set_title("Ridge Regression")
8 l = ax.legend(handles=ax.lines[:len(np.unique(y_train))+1], labels=["Train"])
9 t = g.tick_params(labelsize=12)
10
```



```
In [76]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(rid, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()*(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE: {:.2f}'.format(CV))
```

Cross Validation RMSE: 64.53

## b. Lasso Regression

```
In [77]: 1 las = Lasso(alpha=0.15)
2 las.fit(X=x_train, y=y_train)
3 y_pred_train = las.predict(X=x_train)
4 y_pred_test = las.predict(X=x_test)
5
6 mse_train = MSE(y_train, y_pred_train)
7 mse_test = MSE(y_test, y_pred_test)
8
9 rmse_train = mse_train**(1/2)
10 rmse_test = mse_test**(1/2)
11
12 rmse_dict["lasso"] = {"Train":rmse_train,"Test":rmse_test}
13
14 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
15 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 64.46

Test set RMSE Scaled: 64.26

What would be the RMSE for the majority (80 percentile) of the data?


```
In [78]: 1 th = price_percentile.iloc[80, :].values[0]
2
3 tmpdf = pd.DataFrame({"y_train":y_train, "y_pred_train":y_pred_train})
4 tmpdf = tmpdf[tmpdf["y_train"] <= th]
5
6 tmpdf2 = pd.DataFrame({"y_test":y_test, "y_pred_test":y_pred_test})
7 tmpdf2 = tmpdf2[tmpdf2["y_test"] <= th]
8
9 mse_train = MSE(tmpdf["y_train"].values, tmpdf["y_pred_train"].values)
10 mse_test = MSE(tmpdf2["y_test"].values, tmpdf2["y_pred_test"].values)
11
12 rmse_train = mse_train**(1/2)
13 rmse_test = mse_test**(1/2)
14
15 rmse_m_dict["lasso"] = {"Train":rmse_train,"Test":rmse_test}
16
17 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
18 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 47.75

Test set RMSE Scaled: 47.50



```
In [79]: 1 f, ax = plt.subplots(1,1, figsize=(30, 9), sharex=False)
2 g = sns.pointplot(x=y_train.astype(int), y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test.astype(int), y=y_pred_test, ax=ax, color="red")
4 t = g.set_xlabel("Price")
5 t = g.set_ylabel("Predicted price distribution")
6 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
7 t = g.set_title("Lasso Regression")
8 l = ax.legend(handles=ax.lines[:len(np.unique(y_train))+1], labels=["Train"]
9 t = g.tick_params(labelsize=12)
10
```



```
In [80]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(las, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()*(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE: {:.2f}'.format(CV))
```

Cross Validation RMSE: 64.58

### c. Decision Trees

```
In [81]: 1 dtr = DecisionTreeRegressor(random_state=RANDOM_SEED,max_depth=14,max_featur
2 dtr.fit(X=x_train, y=y_train)
3 y_pred_train = dtr.predict(X=x_train)
4 y_pred_test = dtr.predict(X=x_test)
5
6 mse_train = MSE(y_train, y_pred_train)
7 mse_test = MSE(y_test, y_pred_test)
8
9 rmse_train = mse_train**(1/2)
10 rmse_test = mse_test**(1/2)
11
12 rmse_dict["dt"] = {"Train":rmse_train,"Test":rmse_test}
13
14 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
15 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 58.78

Test set RMSE Scaled: 59.99


What would be the RMSE for the majority (80 percentile) of the data?

```
In [82]: 1 th = price_percentile.iloc[80, :].values[0]
2
3 tmpdf = pd.DataFrame({"y_train":y_train, "y_pred_train":y_pred_train})
4 tmpdf = tmpdf[tmpdf["y_train"] <= th]
5
6 tmpdf2 = pd.DataFrame({"y_test":y_test, "y_pred_test":y_pred_test})
7 tmpdf2 = tmpdf2[tmpdf2["y_test"] <= th]
8
9 mse_train = MSE(tmpdf["y_train"].values, tmpdf["y_pred_train"].values)
10 mse_test = MSE(tmpdf2["y_test"].values, tmpdf2["y_pred_test"].values)
11
12 rmse_train = mse_train**(1/2)
13 rmse_test = mse_test**(1/2)
14
15 rmse_m_dict["dt"] = {"Train":rmse_train,"Test":rmse_test}
16
17 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
18 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 40.82

Test set RMSE Scaled: 41.32

```
In [83]: 1 f, ax = plt.subplots(1,1, figsize=(30, 9), sharex=False)
2 g = sns.pointplot(x=y_train.astype(int), y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test.astype(int), y=y_pred_test, ax=ax, color="red")
4 t = g.set_xlabel("Price")
5 t = g.set_ylabel("Predicted price distribution")
6 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
7 t = g.set_title("Decision Tree Regression")
8 l = ax.legend(handles=ax.lines[:len(np.unique(y_train))+1], labels=["Train"])
9 t = g.tick_params(labelsize=12)
10
```



```
In [85]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(dtr, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()*(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE: {:.2f}'.format(CV))
```

Cross Validation RMSE: 61.04

## d. Random Forests

```
In [86]: 1 rfr = RandomForestRegressor(random_state=RANDOM_SEED, bootstrap=False, crite
2                                     max_depth=20, max_features='sqrt', min_samples_sp
3                                     n_estimators=400)
4 rfr.fit(X=x_train, y=y_train)
5 y_pred_train = rfr.predict(X=x_train)
6 y_pred_test = rfr.predict(X=x_test)
7
8 mse_train = MSE(y_train, y_pred_train)
9 mse_test = MSE(y_test, y_pred_test)
10
11 rmse_train = mse_train**(1/2)
12 rmse_test = mse_test**(1/2)
13
14 rmse_dict["rf"] = {"Train":rmse_train,"Test":rmse_test}
15
16 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
17 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 22.40

Test set RMSE Scaled: 52.83


What would be the RMSE for values less than 250\$?

```
In [87]: 1 th = price_percentile.iloc[80, :].values[0]
2
3 tmpdf = pd.DataFrame({"y_train":y_train, "y_pred_train":y_pred_train})
4 tmpdf = tmpdf[tmpdf["y_train"] <= th]
5
6 tmpdf2 = pd.DataFrame({"y_test":y_test, "y_pred_test":y_pred_test})
7 tmpdf2 = tmpdf2[tmpdf2["y_test"] <= th]
8
9 mse_train = MSE(tmpdf["y_train"].values, tmpdf["y_pred_train"].values)
10 mse_test = MSE(tmpdf2["y_test"].values, tmpdf2["y_pred_test"].values)
11
12 rmse_train = mse_train**(1/2)
13 rmse_test = mse_test**(1/2)
14
15 rmse_m_dict["rf"] = {"Train":rmse_train,"Test":rmse_test}
16
17 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
18 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))
```

Train set RMSE Scaled: 17.18

Test set RMSE Scaled: 36.84

```
In [88]: 1 f, ax = plt.subplots(1,1, figsize=(30, 9), sharex=False)
2 g = sns.pointplot(x=y_train.astype(int), y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test.astype(int), y=y_pred_test, ax=ax, color="red")
4 t = g.set_xlabel("Price")
5 t = g.set_ylabel("Predicted price distribution")
6 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
7 t = g.set_title("Random Forest Regression")
8 l = ax.legend(handles=ax.lines[:len(np.unique(y_train))+1], labels=["Train"]
9 t = g.tick_params(labelsize=12)
10
```



```
In [89]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(rfr, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()*(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE: {:.2f}'.format(CV))
```


Cross Validation RMSE: 53.17

## Feature importance

```

In [90]: 1 x_final = listings.drop(['price'], axis=1)
2
3 imp = pd.DataFrame({'Input Variable': x_final.columns, 'Importance (%)': -np
4 top20 = imp.iloc[:20,:]}
5
6 f, ax = plt.subplots(1,1,figsize=(20,9))
7 g = sns.barplot(x="Importance (%)", y="Input Variable", data=top20, ax=ax)
8 t = g.set_xlabel("Importance (%)")
9 t = g.set_ylabel("Input Variable")
10 t = g.set_title("Feature Importance")

```



## e. Gradient Boosting regression

```

In [91]: 1 gbr = GradientBoostingRegressor(random_state=RANDOM_SEED,
2                                     n_estimators=400,
3                                     subsample=1,
4                                     criterion="mse",
5                                     max_depth=5,
6                                     max_features="sqrt",
7                                     min_samples_split=35,
8                                     min_samples_leaf=5,
9                                     learning_rate=0.1)
10 gbr.fit(X=x_train, y=y_train)
11 y_pred_train = gbr.predict(X=x_train)
12 y_pred_test = gbr.predict(X=x_test)
13
14 mse_train = MSE(y_train, y_pred_train)
15 mse_test = MSE(y_test, y_pred_test)
16
17 rmse_train = mse_train**(1/2)
18 rmse_test = mse_test**(1/2)
19
20 rmse_dict["gbr"] = {"Train":rmse_train,"Test":rmse_test}
21
22 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
23 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))

```

Train set RMSE Scaled: 46.30

Test set RMSE Scaled: 53.15

What would be the RMSE for values less than 250\$?

```


In [92]: 1 th = price_percentile.iloc[80, :].values[0]
2
3 tmpdf = pd.DataFrame({"y_train":y_train, "y_pred_train":y_pred_train})
4 tmpdf = tmpdf[tmpdf["y_train"] <= th]
5
6 tmpdf2 = pd.DataFrame({"y_test":y_test, "y_pred_test":y_pred_test})
7 tmpdf2 = tmpdf2[tmpdf2["y_test"] <= th]
8
9 mse_train = MSE(tmpdf["y_train"].values, tmpdf["y_pred_train"].values)
10 mse_test = MSE(tmpdf2["y_test"].values, tmpdf2["y_pred_test"].values)
11
12 rmse_train = mse_train**(1/2)
13 rmse_test = mse_test**(1/2)
14
15 rmse_m_dict["gbr"] = {"Train":rmse_train,"Test":rmse_test}
16
17 print("Train set RMSE Scaled: {:.2f}".format(rmse_train))
18 print("Test set RMSE Scaled: {:.2f}".format(rmse_test))

```

Train set RMSE Scaled: 33.64


Test set RMSE Scaled: 37.08

```
In [93]: 1 f, ax = plt.subplots(1,1, figsize=(30, 9), sharex=False)
2 g = sns.pointplot(x=y_train.astype(int), y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test.astype(int), y=y_pred_test, ax=ax, color="red")
4 t = g.set_xlabel("Price")
5 t = g.set_ylabel("Predicted price distribution")
6 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
7 t = g.set_title("Gradient Boosting Regression")
8 l = ax.legend(handles=ax.lines[:,len(np.unique(y_train))+1], labels=["Train"])
9 t = g.tick_params(labelsize=12)
10
```



## Feature importance

```
In [94]: 1 x_final = listings.drop(['price'], axis=1)
2
3 imp = pd.DataFrame({'Input Variable': x_final.columns, 'Importance (%)': -np
4 top20 = imp.iloc[:20,:]}
5
6 f, ax = plt.subplots(1,1,figsize=(20,9))
7 g = sns.barplot(x="Importance (%)", y="Input Variable", data=top20, ax=ax)
8 t = g.set_xlabel("Importance (%)")
9 t = g.set_ylabel("Input Variable")
10 t = g.set_title("Feature Importance")
```





```
In [95]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(gbr, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()**(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE: {:.2f}'.format(CV))
```

Cross Validation RMSE: 52.94

Let's plot a quick bar chart that captures all this information

```
In [96]: 1 rmse_df = pd.DataFrame(rmse_dict)
2 rmse_df = rmse_df.T.reset_index(drop=False)
3 rmse_df.columns = ["Regressor", "Test RMSE", "Train RMSE"]
4 rmse_df
```

Out[96]:

	Regressor	Test RMSE	Train RMSE
0	Ridge	65.168006	64.395215
1	lasso	64.256259	64.462710
2	dt	59.988126	58.779791
3	rf	52.831975	22.402276
4	gbr	53.151973	46.301878

```
In [97]: 1 rmse_m_df = pd.DataFrame(rmse_m_dict)
2 rmse_m_df = rmse_m_df.T.reset_index(drop=False)
3 rmse_m_df.columns = ["Regressor", "Test RMSE", "Train RMSE"]
4 rmse_m_df
```

Out[97]:


	Regressor	Test RMSE	Train RMSE
0	Ridge	49.447526	47.920813
1	lasso	47.503708	47.751292
2	dt	41.324700	40.822763
3	rf	36.838510	17.179769
4	gbr	37.078385	33.636091

```

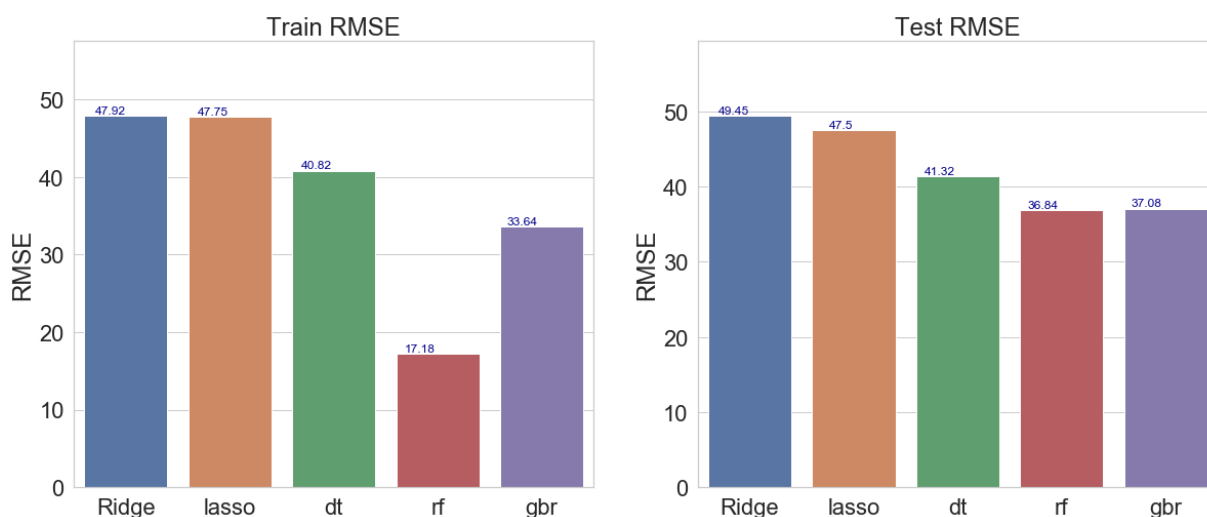
In [98]: 1 def plot_bar(data, x, y, ax, hue=None, title="", xlabel="", ylabel="",
2         xrot=0, yrot=0, highlight_max_min=True,
3         plot_percentiles=[], plot_mean=True,
4         point_plot=False, annot=True, legend=False):
5     if highlight_max_min:
6         clr = []
7         for v in data[y].values:
8             if v < data[y].max():
9                 if v > data[y].min():
10                    clr.append('lightblue')
11                else:
12                    clr.append('lightgreen')
13            else:
14                clr.append('darksalmon')
15        g = sns.barplot(x=x, y=y, data=data, ax=ax, palette=clr)
16    else:
17        g = sns.barplot(x=x, y=y, data=data, ax=ax, hue=hue)
18
19    if len(plot_percentiles) > 0:
20        for p in plot_percentiles:
21            v = np.percentile(data[y].values, p)
22            plt.axhline(v, 1, 0, color='grey').set_linestyle("--")
23
24    if plot_mean:
25        v = data[y].mean()
26        plt.axhline(v, 1, 0, color='k').set_linestyle("--")
27
28    if point_plot:
29        g1 = sns.pointplot(x=x, y=y, data=data, ax=ax, color="darkblue")
30    if xrot != 0:
31        g.set_xticklabels(rotation=xrot, labels=g.get_xticklabels())
32    if yrot != 0:
33        g.set_yticklabels(rotation=yrot, labels=g.get_yticklabels())
34    if annot:
35        # Add labels to the plot
36        style = dict(size=12, color='darkblue')
37        s1 = data[y].values
38        counter = 0
39        for idx, row in data.iterrows():
40            rx, ry = row[x], row[y]
41            if type('str') == type(idx):
42                ax.text(counter, ry, str(np.round(ry, 2)),
43                        **style, va="bottom", ha='right')
44            else:
45                ax.text(idx*0.99, ry, str(np.round(s1[idx], 2)),
46                        **style, va="bottom", ha='right')
47            counter += 1
48    g.set(xlabel=xlabel, ylabel=ylabel, title=title)
49    ax.set_ylim([0, data[y].max() * 1.2])
50    if legend:
51        ax.legend(handles=ax.lines[:len(data) + 1], labels=[y])
52

```

```
In [104]: 1 f, ax = plt.subplots(1,2,figsize=(20,8), sharey=False)
2 # g = sns.barplot(x="Regressor", y="Train RMSE", data=rmse_df, ax=ax[0])
3 # t = g.set(title="Train RMSE", ylabel="RMSE", xlabel="")
4 # g = sns.barplot(x="Regressor", y="Test RMSE", data=rmse_df, ax=ax[1])
5 # t = g.set(title="Test RMSE", ylabel="RMSE", xlabel="")
6
7 g = plot_bar(x="Regressor", y="Train RMSE", data=rmse_df, ax=ax[0], plot_mean
8 g = plot_bar(x="Regressor", y="Test RMSE", data=rmse_df, ax=ax[1], plot_mean
9
```



```
In [105]: 1 f, ax = plt.subplots(1,2,figsize=(20,8), sharey=False)
2 # g = sns.barplot(x="Regressor", y="Train RMSE", data=rmse_m_df, ax=ax[0])
3 # t = g.set(title="Train RMSE", ylabel="RMSE", xlabel="")
4 # g = sns.barplot(x="Regressor", y="Test RMSE", data=rmse_m_df, ax=ax[1])
5 # t = g.set(title="Test RMSE", ylabel="RMSE", xlabel="")
6
7 g = plot_bar(x="Regressor", y="Train RMSE", data=rmse_m_df, ax=ax[0], plot_m
8 g = plot_bar(x="Regressor", y="Test RMSE", data=rmse_m_df, ax=ax[1], plot_me
```



The overall RMSE values, and the general "fit" of the model onto the data has substantially improved.

It can also be noted that the noise level increases significantly in later price bins, albeit lesser than before we used SMOTE to oversample values. The lower price bins are much more consistent with lesser noise than the higher ones.

Further exploration can be conducted to improve the fit of the regressors onto the dataset, and get lower RMSEs.