

AirBnb listings file modeling

```
In [1]: 1 # Put these at the top of every notebook, to get automatic reloading and inl
2 from IPython.core.display import display, HTML
3 import pandas as pd
4 import warnings
5 import ast
6 warnings.filterwarnings('ignore')
7
8 %reload_ext autoreload
9 %autoreload 1
10 %matplotlib inline
11
12 pd.set_option('display.max_rows', 500)
13 pd.set_option('display.max_columns', 500)
14 pd.set_option('display.width', 1000)
15
16 display(HTML("<style>.container { width:100% !important; }</style>"))
```

```
In [1]: 1 import os
2 import seaborn as sns
3 import pandas as pd
4 import math
5
6 import sklearn.model_selection as cv
7
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.model_selection import train_test_split
10 from sklearn.decomposition import PCA
11 from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor, GradientBoostingRegressor
12 from sklearn.model_selection import GridSearchCV
13
14 from sklearn.metrics import mean_squared_error as MSE
15
16 from imblearn.over_sampling import SMOTE
17
18 from Utils.UtilsGeoViz import *
19 from Utils.UtilsViz import *
20 from Utils.DataUtils import *
21
22 RANDOM_SEED = 42
```

```
In [3]: 1 data_path = os.path.join(os.getcwd(), "../data/cleaned_listings.csv")
2 listings = pd.read_csv(data_path, index_col="id")
3 display(listings.shape)
```


(48855, 65)

Plot the distribution

Let's plot the percentile for price

```
In [4]: 1 percentiles = list(range(0,101, 1))
        2 price_percentile = {}
        3 for p in percentiles:
        4     price_percentile[p] = np.percentile(listings['price'].values, p)
        5
        6 sns.set(style="white")
        7 price_percentile = pd.DataFrame.from_dict(price_percentile, orient='index')
        8 price_percentile.plot(kind='bar', figsize=(20,9), grid=True, legend=False)
        9 plt.title("Percentile distribution of price")
       10 plt.xlabel("Percentile")
       11 plt.ylabel("Price")
```

Out[4]: Text(0, 0.5, 'Price')




```

In [5]: 1 f, ax = plt.subplots(1,1,figsize=(20,9))
2 g = sns.countplot(x="price", data=listings, ax=ax)
3 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
4 t = g.set_title("price distribution")
5 median_idx = np.where(np.sort(listings["price"].unique())==listings["price"]
6 plt.axvline(x=median_idx, linestyle="-.")
7 percentile_75_idx = np.where(np.sort(listings["price"].unique())==price_perc
8 plt.axvline(x=percentile_75_idx, linestyle="--")
9 percentile_90_idx = np.where(np.sort(listings["price"].unique())==price_perc
10 plt.axvline(x=percentile_90_idx, linestyle=":")

```

Out[5]: <matplotlib.lines.Line2D at 0x12b4d870>



Quick helper functions

```

In [6]: 1 def roundto(row, base=5):
2         return int(base * round(float(row) / base))
3
4 # Get the index of the price columns
5 def get_index(vallist, val):
6     return vallist.index(val)

```

1. Oversampling using SMOTE

```
In [7]: 1 def check_rep(row):
2         if (row <= 200) | (row==250) | (row==350) | (row==450) | (row==550) :
3             return 0
4         elif (row > 200) & (row < 300) & (row != 250):
5             return 1
6         elif (row > 300) & (row < 400) & (row != 350):
7             return 2
8         else:
9             return 3
10
11 listings["flag_ur"] = listings["price"].apply(check_rep)
```

```
In [8]: 1 vcs = listings["flag_ur"].value_counts()
2         vcs
```

```
Out[8]: 0    43321
1      3093
3      1624
2       817
Name: flag_ur, dtype: int64
```

```
In [9]: 1 ycol = ["flag_ur"]
2         xcol = [i for i in listings.columns if i not in ycol]
3
4         x = listings[xcol].values
5         y = listings[ycol].values
6
7         smote_sampling_strategy = {
8             1: int(vcs[1]*2)
9             ,2: int(vcs[2]*2)
10            ,3: int(vcs[3]*2)
11        }
12         sm = SMOTE(random_state=RANDOM_SEED, sampling_strategy=smote_sampling_strate
13         # Fit the smote onto the sample
14         x_new, y_new = sm.fit_sample(x, y)
15
16         # Drop the flag column
17         listings.drop(labels=["flag_ur"], axis=1, inplace=True)
18
19         # -----
20         # Overwrite X and Y
21         price_index = get_index(list(listings.columns), "price")
22
23         y = x_new[:, price_index]
24         x = np.delete(x_new, price_index, axis=1)
25         for i in range(len(y)):
26             y[i] = roundto(y[i])
```


c:\users\sriharis\appdata\local\programs\python\python37-32\lib\site-packages\s
klearn\utils\validation.py:761: DataConversionWarning: A column-vector y was pa
ssed when a 1d array was expected. Please change the shape of y to (n_samples,
,), for example using ravel().
y = column_or_1d(y, warn=True)

```
In [10]: 1 print(
2         " Old size :", listings.shape, "\n",
3         "New size :", x.shape
4     )
```

Old size : (48855, 65)

New size : (54389, 64)

```
In [11]: 1 f, ax = plt.subplots(1, 2, figsize=(30, 9), sharey=True)
2 g1 = sns.distplot(listings["price"], ax=ax[0], kde=False)
3 t = g1.set_title("Original price distribution")
4
5 g2 = sns.distplot(y, ax=ax[1], kde=False)
6 t = g2.set_title("SMOTE price distribution")
7
```



2. Transformation

```
In [12]: 1 x_cols = listings.drop(['price'], axis=1)
2 X = pd.DataFrame(data=x, columns=x_cols.columns)
```

```
In [13]: 1 X_log = X.copy()
2 X_sqr = X.copy()
3 X_sqrt = X.copy()
```

```
In [14]: 1 # Taking log, square and square root transformations of x
2 for i in range(X_log.shape[1]):
3     X_log.iloc[:,i] = np.log(X_log.iloc[:,i] + 1)
4     X_sqr.iloc[:,i] = np.square(X_sqr.iloc[:,i])
5     X_sqrt.iloc[:,i] = np.sqrt(X_sqrt.iloc[:,i])
6 X_log.columns = X_log.columns.map(lambda x: x + '_log')
7 X_sqr.columns = X_sqr.columns.map(lambda x: x + '_sqr')
8 X_sqrt.columns = X_sqrt.columns.map(lambda x: x + '_sqrt')
```

```
In [15]: 1 X_sqrt.shape
```

Out[15]: (54389, 64)

```
In [16]: 1 # Appending X, X_log, X_sqr and X_sqrt
2 X_final = pd.concat([X, X_log, X_sqr, X_sqrt], axis=1)
3 X_final.shape
```

Out[16]: (54389, 256)

Prediction

```
In [17]: 1 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.30, ra
```

```
In [18]: 1 rfr = RandomForestRegressor(random_state=RANDOM_SEED)
2 rfr.fit(X=x_train, y=y_train)
3 y_pred_train = rfr.predict(X=x_train)
4 y_pred_test = rfr.predict(X=x_test)
5
6 mse_train = MSE(y_train, y_pred_train)
7 mse_test = MSE(y_test, y_pred_test)
8
9 rmse_train = mse_train**(1/2)
10 rmse_test = mse_test**(1/2)
11
12 print("Train set RMSE with Transformation: {:.2f}".format(rmse_train))
13 print("Test set RMSE with Transformation: {:.2f}".format(rmse_test))
```

c:\users\sriharis\appdata\local\programs\python\python37-32\lib\site-packages\s
klearn\ensemble\forest.py:246: FutureWarning: The default value of n_estimators
will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

Train set RMSE with Transformation: 24.25

Test set RMSE with Transformation: 56.60

```
In [19]: 1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(rfr, x_train, y_train, scoring='neg_mean_sq
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()**(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE with Transformation: {:.2f}'.format(CV))
```

Cross Validation RMSE with Transformation: 57.38

3. Train test split

```
In [20]: 1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, ra
```

4. Standardisation

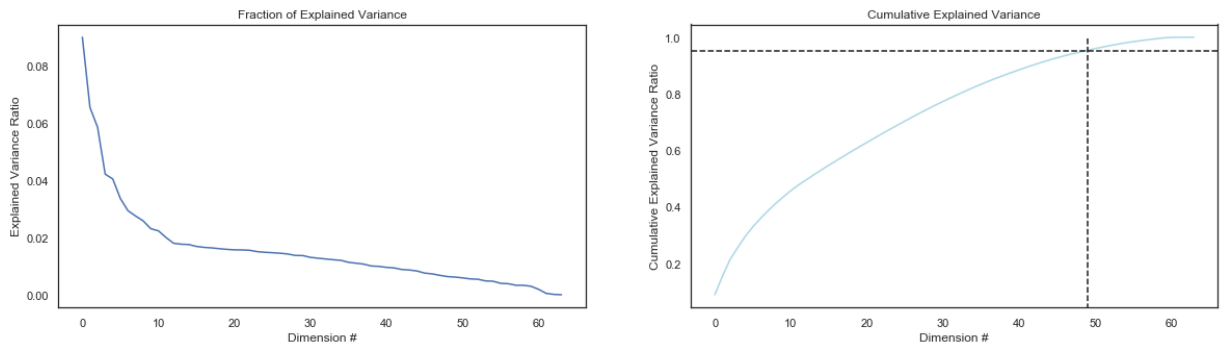
```
In [21]: 1 standard_scaler = StandardScaler()
2 x_train = standard_scaler.fit_transform(x_train)
3 x_test = standard_scaler.transform(x_test)
```

5. PCA

```
In [22]: 1 pca_naive = PCA(n_components=None)
2 pca_naive.fit(x_train)
3 f, ax = plt.subplots(1,2,figsize=(20,5))
4 plot_naive_variance(pca_naive, ax[0])
5 num_pc = plot_pca_var_cum(pca_naive, ax=ax[1], cutoff=0.95)
6 print("Number of components : ", num_pc)
```

c:\users\sriharis\appdata\local\programs\python\python37-32\lib\site-packages\matplotlib\book__init__.py:1725: UserWarning: Saw kwargs ['c', 'color'] which are all aliases for 'color'. Kept value from 'color' seen=seen, canon=canonical, used=seen[-1]))

Number of components : 49



```
In [23]: 1 # pca = PCA(n_components=num_pc)
2 pca = PCA(n_components=None)
3 pca.fit(x_train)
4 x_train_pca = pca.transform(x_train)
5 x_test_pca = pca.transform(x_test)
```

Prediction

Random Forest Regressor

With PCA

In [24]:

```

1 rfr = RandomForestRegressor(random_state=RANDOM_SEED)
2 rfr.fit(X=x_train_pca, y=y_train)
3 y_pred_train = rfr.predict(X=x_train_pca)
4 y_pred_test = rfr.predict(X=x_test_pca)
5
6 mse_train = MSE(y_train, y_pred_train)
7 mse_test = MSE(y_test, y_pred_test)
8
9 rmse_train = mse_train**(1/2)
10 rmse_test = mse_test**(1/2)
11
12 print("Train set RMSE with PCA: {:.2f}".format(rmse_train))
13 print("Test set RMSE with PCA: {:.2f}".format(rmse_test))

```

c:\users\sriharis\appdata\local\programs\python\python37-32\lib\site-packages\s
klearn\ensemble\forest.py:246: FutureWarning: The default value of n_estimators
will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

Train set RMSE with PCA: 26.72

Test set RMSE with PCA: 62.94

In [25]:

```

1 # Use cross-validation on Train data
2
3 CV_scores = - cv.cross_val_score(rfr, x_train_pca, y_train, scoring='neg_mea
4
5 # Compute the 10-folds CV
6 CV = CV_scores.mean()**(1/2)
7
8 # Print Train CV accuracy
9 print('Cross Validation RMSE with PCA: {:.2f}'.format(CV))

```

Cross Validation RMSE with PCA: 63.11

In [26]:

```

1 f, ax = plt.subplots(1,1, figsize=(30, 9))
2 g = sns.pointplot(x=y_train, y=y_pred_train, ax=ax)
3 g = sns.pointplot(x=y_test, y=y_pred_test, ax=ax, color="red")
4
5 t = g.set_xlabel("Price")
6 t = g.set_ylabel("Predicted price distribution")
7 t = g.set_xticklabels(g.get_xticklabels(), rotation=90)
8 t = g.set_title("PCA")

```

